

Knowledge Graphs

Few slides from CS224W @ Stanford

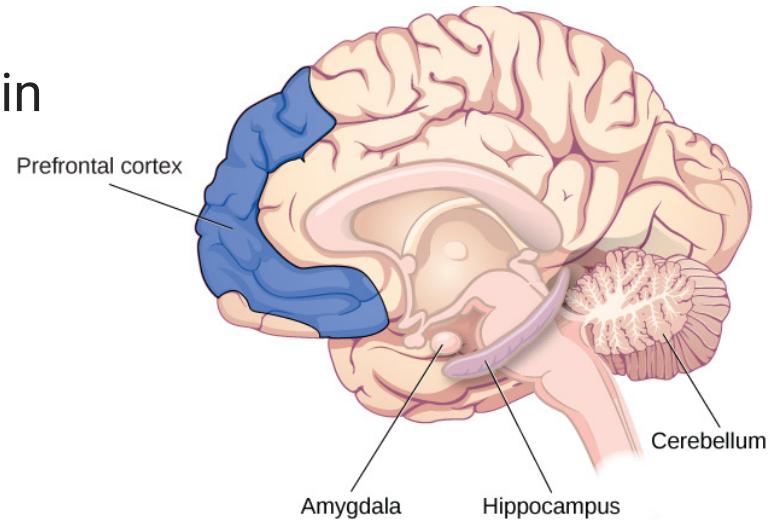
Mrinmaya Sachan
msachan@ethz.ch

The need for a Knowledge Store in AI systems

There is increasing evidence that humans **store information** in a part of their brains (most probably the **hippocampus**).

They can readily **retrieve** information as needed:

- When a memory is **created**, information flows from the **cortex**, the part of the brain rich in nerve cells, to the **hippocampus**, the central switching point for memories in the brain.
- The information flows in the opposite direction when we **retrieve** a memory



Perhaps, AI models should be able to store and interact with information/knowledge as well.

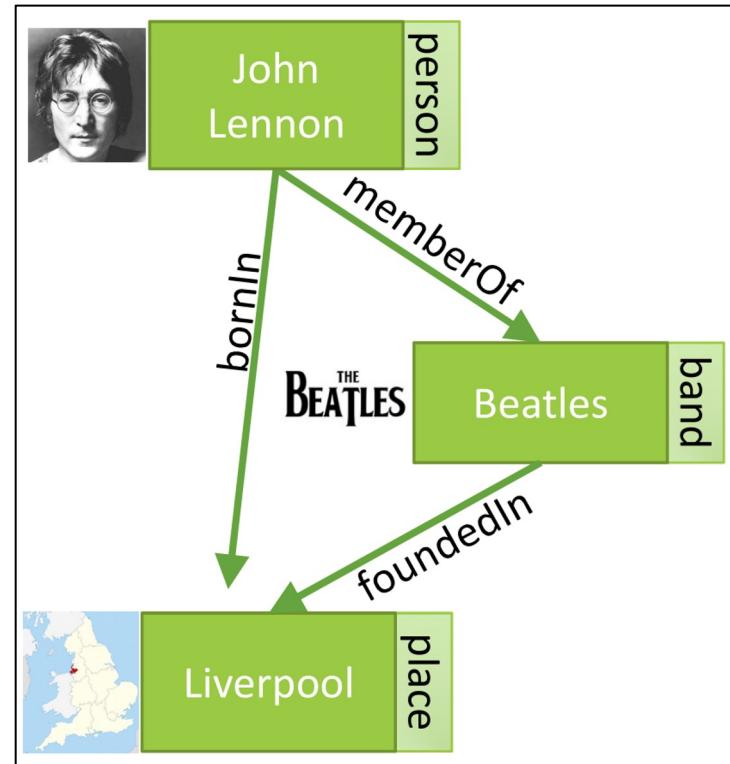
Knowledge Graphs

Describes **entities** in the real world and their **relationships**

Nodes are **entities**

Nodes are labeled with their **types**

Edges between two nodes capture
relationships between entities



Easy and sound way of representing world knowledge in a way that is usable by computers

Knowledge Graphs: Another representation

Represented as set of **Subject, Relation, Object** triples

- (John Lennon, memberOf, Beatles)
- (Beatles, foundedIn, Liverpool)
- (John Lennon, bornin, Liverpool)

Industry Scale Knowledge Graphs

GOOGLE KNOWLEDGE GRAPH

- <https://developers.google.com/knowledge-graph/>

AMAZON PRODUCT GRAPH

FACEBOOK GRAPH API

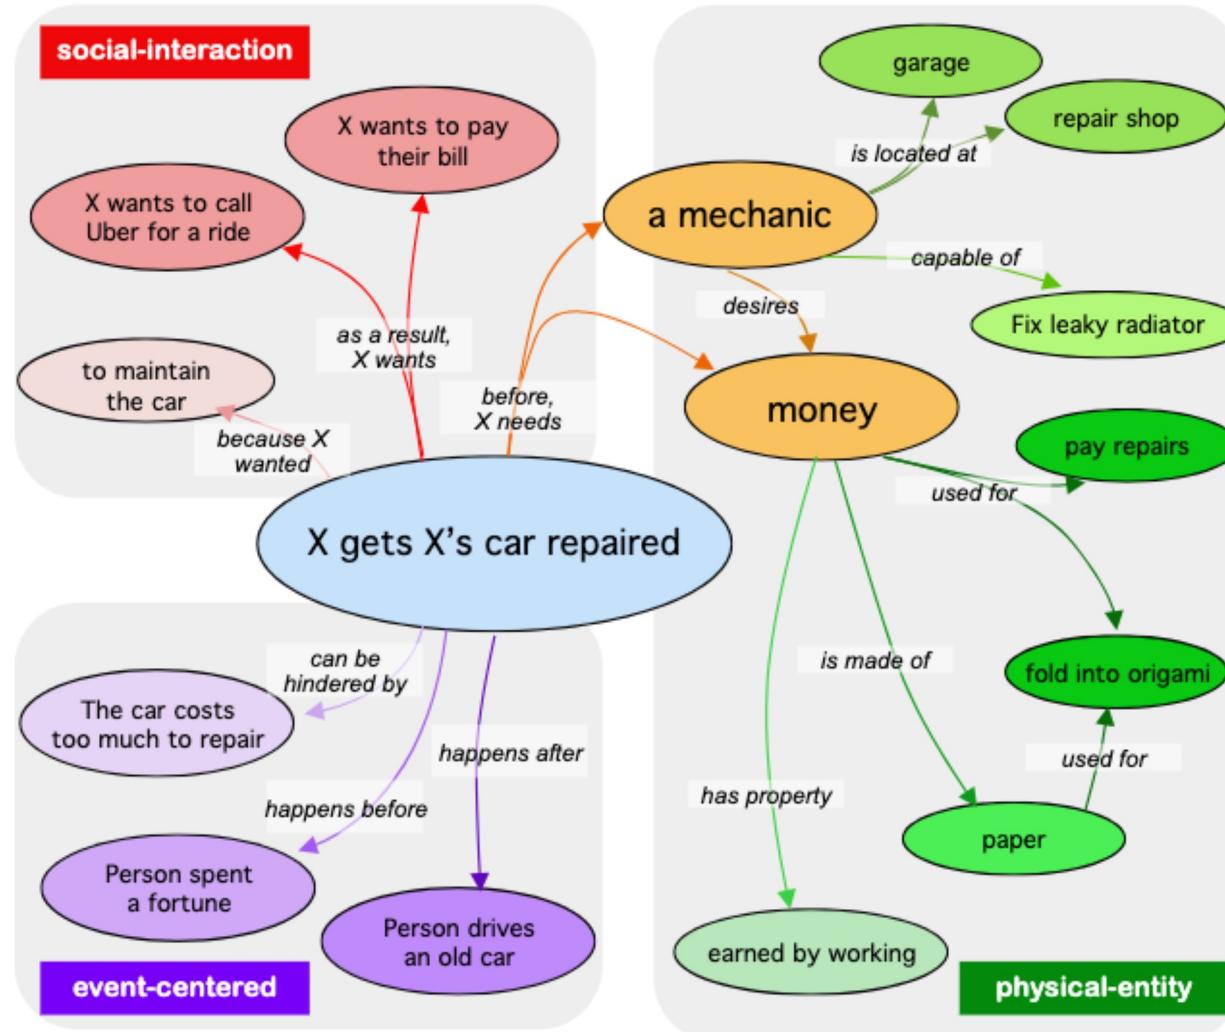
- <https://developers.facebook.com/docs/graph-api/>

MICROSOFT SATORI

LINKEDIN KNOWLEDGE GRAPH

Academic Knowledge Graphs

YAGO (MPI), ConcepNet (MIT Media Labs), NELL (CMU), ATOMIC (UW/AI2)

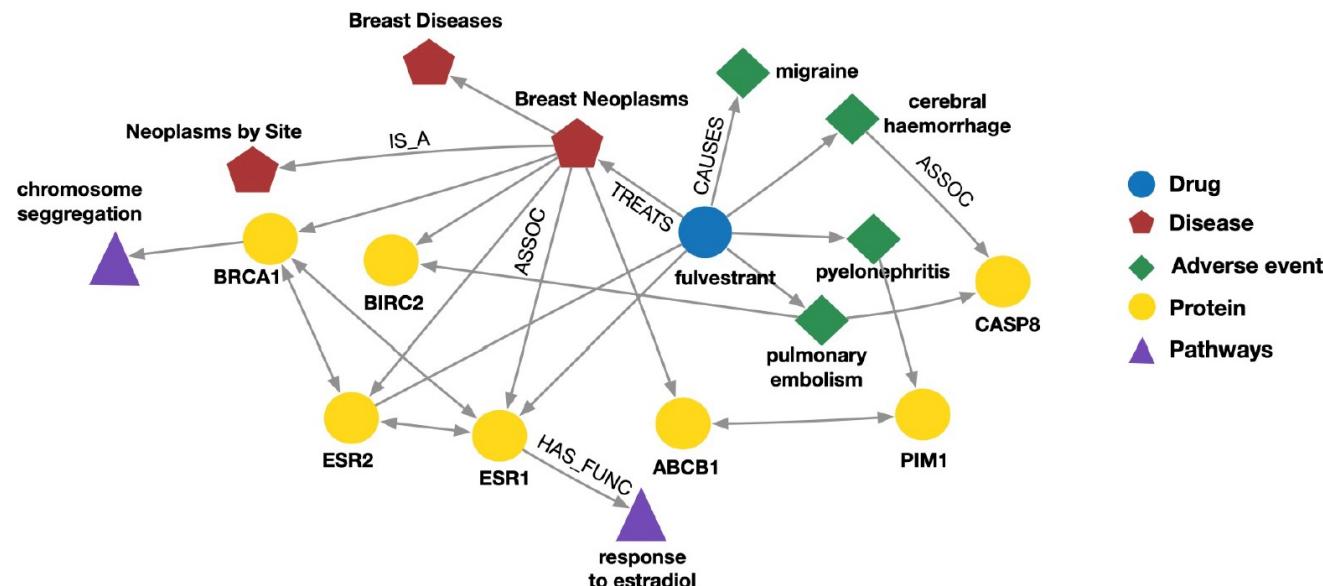


Domain specific KGs: Biology and Healthcare

MeSH (Medline/PubMed), UMLS (Unified Medical Language System),
Disease Ontology, Cell Ontology, Gene Ontology

Node types: drug, disease, adverse event, protein, pathways

Relation types: has_func, causes, assoc, treats, is_a



Applications of Knowledge Graphs

Serving information:

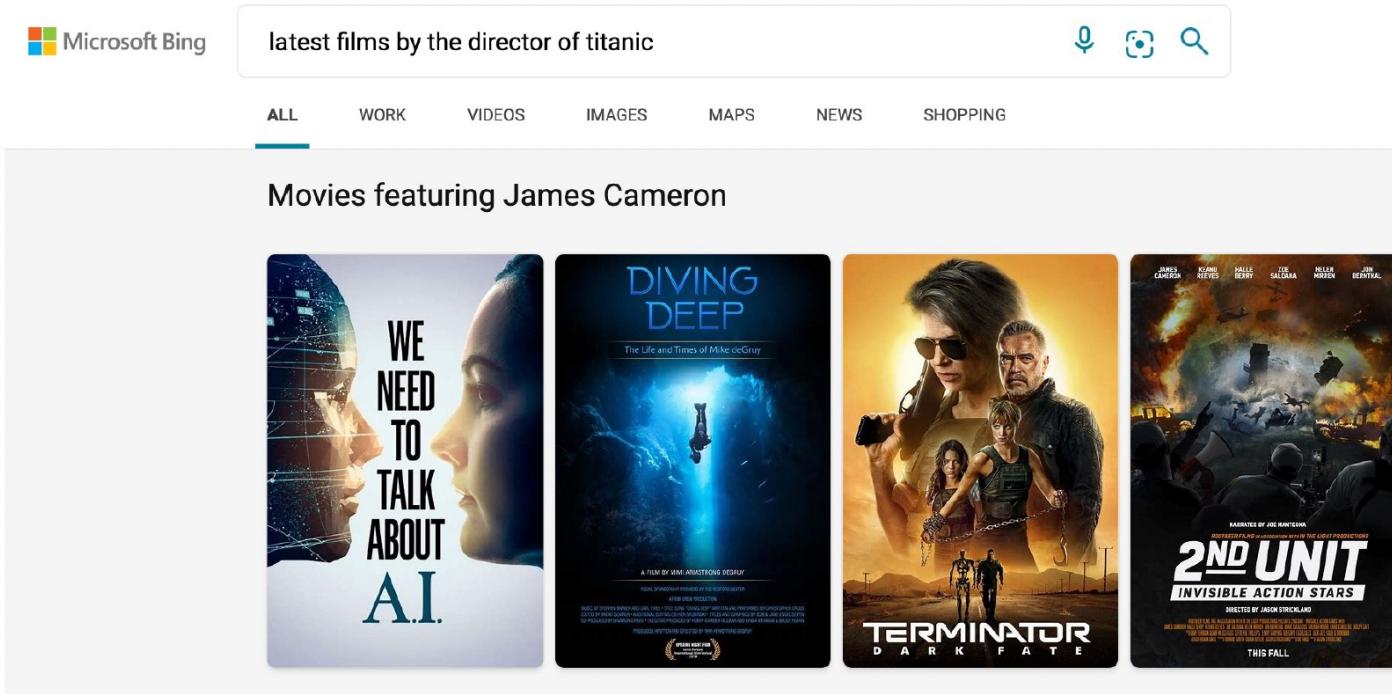


Image credit: Bing

Question answering and conversation agents

However...

- KGs are inherently **symbolic**
 - We can search a KG and run interpreted queries such as “`foundedIn(Beatles, ?)`”
- It is challenging to incorporate them in deep learning architectures that require **continuous representations**

Thus...

- We need a way to embed a knowledge graph into a **continuous** vector space!

Knowledge Graph Representations

Knowledge Graph Representations

Given a collection of facts denoted as (h, r, t) from our KG

We wish to learn embeddings h, r, t for the head and tail entities and relation. Let $h, t \in \mathbb{R}^d$ and $r \in \mathbb{R}^k$

How can we learn these representations?

The key idea is to learn an embedding space such that h when “combined” with r is close to t

Knowledge Graph Representations

The key idea is to learn an embedding space such that \mathbf{h} when “combined” with \mathbf{r} is close to \mathbf{t}

Let $f_r(\mathbf{h}, \mathbf{t})$ measure the **compatibility** of the embedding of the head entity with the embedding of the tail entity given the embedding of the relation.

- In other words, f quantifies the **plausibility** of the embedded representation of a given fact.

We want:

Score f_r should be high if (h, r, t) exists in the KB, else f is low

A more refined way to think about this...

Objective:

$$\min \sum_{(h, r, t) \in KG} L(f_\theta(h, r), t)$$

Here, f_θ combines h and r

L measures the closeness between “the combination of h and r ” and t

Twist: To avoid trivial solutions, minimize L for known facts and maximize L for unknown facts generated using **negative sampling**

$$\min \sum_{(h, r, t) \in KG} L(f_\theta(h, r), t) - \sum_{(h, r, t^-) \notin KG} L(f_\theta(h, r), t^-)$$

Some popular KG embedding methods

Different forms of $L(f_\theta(\mathbf{h}, \mathbf{r}), \mathbf{t})$. Let $\mathbf{e}_i = \mathbf{h}$ and $\mathbf{e}_j = \mathbf{t}$

SE (Bordes et al., 2011)	$\ \mathbf{W}_r^{(L)}\mathbf{e}_i - \mathbf{W}_r^{(R)}\mathbf{e}_j\ _p$
NTN (Socher et al., 2013)	$\mathbf{u}_r^T f \left(\mathbf{e}_i \mathbf{W}_r^{[1\dots k]} \mathbf{e}_j + \mathbf{V}_r \begin{bmatrix} \mathbf{e}_i \\ \mathbf{e}_j \end{bmatrix} + \mathbf{b}_r \right)$
TransE (Bordes et al., 2013)	$\ \mathbf{e}_i + \mathbf{r} - \mathbf{e}_j\ _p$
TransH (Wang et al., 2014)	$\ (\mathbf{e}_i - \mathbf{w}_r^T \mathbf{e}_i \mathbf{w}_r) + \mathbf{d}_r - (\mathbf{e}_j - \mathbf{w}_r^T \mathbf{e}_j \mathbf{w}_r)\ _2^2$
TransR (Lin et al., 2015)	$\ \mathbf{e}_i \mathbf{W}_r + \mathbf{r} - \mathbf{e}_j \mathbf{W}_r\ _2^2$
TransD (Ji et al., 2015)	$- \ (\mathbf{r}_p \mathbf{e}_i p^T + \mathbf{I}) \mathbf{e}_i + \mathbf{r} - (\mathbf{r}_p \mathbf{e}_j p^T + \mathbf{I}) \mathbf{e}_j\ _2^2$
DistMult (Yang et al., 2014)	$\langle \mathbf{e}_i, \mathbf{r}, \mathbf{e}_j \rangle$
ComplEx (Trouillon et al., 2016)	$\text{Re}(\langle \mathbf{e}_i, \mathbf{r}, \bar{\mathbf{e}}_j \rangle)$
HolE (Nickel et al., 2016)	$\langle \mathbf{r}, \mathbf{e}_i \otimes \mathbf{e}_j \rangle$
SimpleE (Kazemi and Poole, 2018)	$\frac{1}{2} (\langle \mathbf{e}_i^{(h)}, \mathbf{r}, \mathbf{e}_j^{(t)} \rangle + \langle \mathbf{e}_j^{(h)}, \mathbf{r}^{(inv)}, \mathbf{e}_i^{(t)} \rangle)$
ConvE (Dettmers et al., 2018)	$\langle \sigma(vec(\sigma([\mathbf{e}_i; \mathbf{r}] \circ \omega)) \mathbf{W}), \mathbf{e}_j \rangle$
RotatE (Sun et al., 2019)	$- \ \mathbf{e}_i \bullet \mathbf{r} - \mathbf{e}_j\ ^2$
HypER (Balažević et al., 2019a)	$\langle \sigma(vec(\sigma(\mathbf{e}_i * vec^{-1}(\mathbf{w}_r \mathbf{H}))) \mathbf{W}), \mathbf{e}_j \rangle$
TuckER (Balažević et al., 2019b)	$\mathcal{W} \times_1 \mathbf{e}_i \times_2 \mathbf{w}_r \times_3 \mathbf{e}_j$

1) TransE

(Bordes et al. 2013)

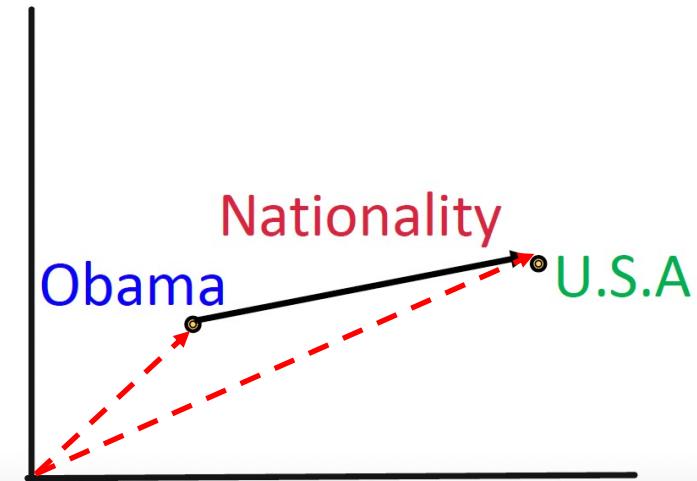
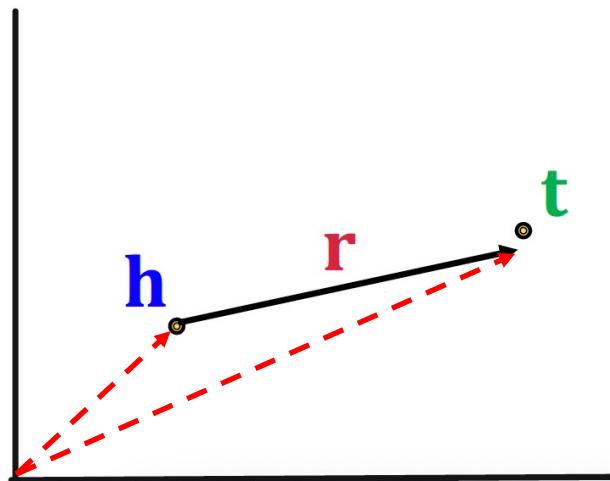
■ Intuition: Translation

For a triple (h, r, t) , let $\mathbf{h}, \mathbf{r}, \mathbf{t} \in \mathbb{R}^d$
be embedding vectors.

embedding vectors
will appear in
boldface

■ TransE: $\mathbf{h} + \mathbf{r} \approx \mathbf{t}$ if the given link exists else $\mathbf{h} + \mathbf{r} \neq \mathbf{t}$

Entity scoring function: $f_r(h, t) = -||\mathbf{h} + \mathbf{r} - \mathbf{t}||$



Algorithm 1 Learning TransE

input Training set $S = \{(h, r, t)\}$, entities and rel. sets E and R , margin γ , embeddings dim. k .

- 1: **initialize** $r \leftarrow \text{uniform}(-\frac{6}{\sqrt{k}}, \frac{6}{\sqrt{k}})$ for each $r \in R$
- 2: $r \leftarrow r / \|r\|$ for each $r \in R$
- 3: $e \leftarrow \text{uniform}(-\frac{6}{\sqrt{k}}, \frac{6}{\sqrt{k}})$ for each entity $e \in E$
- 4: **loop**
- 5: $e \leftarrow e / \|e\|$ for each entity $e \in E$
- 6: $S_{batch} \leftarrow \text{sample}(S, b)$ // sample a minibatch of size b
- 7: $T_{batch} \leftarrow \emptyset$ // initialize the set of pairs of triplets
- 8: **for** $(h, r, t) \in S_{batch}$ **do**
- 9: $(h', r, t') \leftarrow \text{sample}(S'_{(h, r, t)})$ // sample a corrupted triplet
- 10: $T_{batch} \leftarrow T_{batch} \cup \{((h, r, t), (h', r, t'))\}$
- 11: **end for**
- 12: Update embeddings w.r.t.

$$\sum_{((h, r, t), (h', r, t')) \in T_{batch}} \nabla [\gamma + d(\mathbf{h} + r, \mathbf{t}) - d(\mathbf{h}' + r, \mathbf{t}')]_+$$

positive sample negative sample

13: **end loop**

Contrastive loss: Favors lower distance (or higher score) for valid triplets, high distance (or lower score) for corrupted ones

Initialize relations r and entities e uniformly, then normalize.
 γ is the margin.

Sample triplet (h', r, t) that does not appear in the KG.

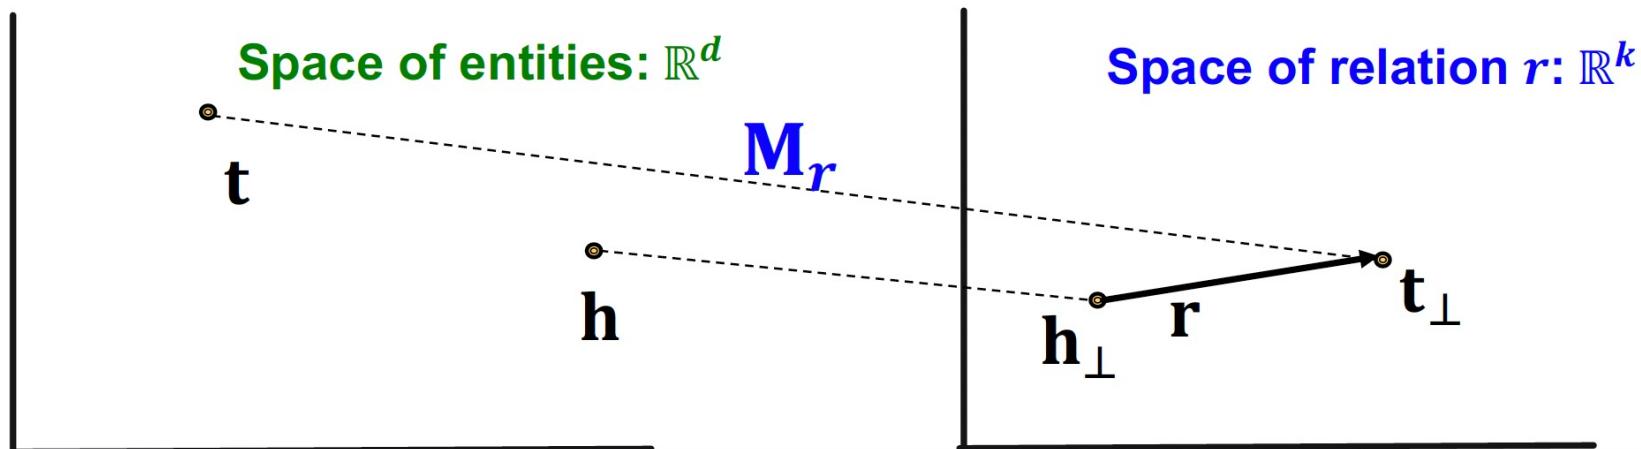
d represents distance
(negative of score)

2) TransR

(Lin et al. 2015)

- **TransE** models translation of any relation in the **same** embedding space.
- Can we design a new space for each relation and do translation in **relation-specific space**?

- **TransR**: model **entities** as vectors in the entity space \mathbb{R}^d and model each **relation** as vector in relation space $\mathbf{r} \in \mathbb{R}^k$ with $\mathbf{M}_r \in \mathbb{R}^{k \times d}$ as the **projection matrix**.
Use \mathbf{M}_r to **project** from entity space \mathbb{R}^d to relation space \mathbb{R}^k !
- $\mathbf{h}_\perp = \mathbf{M}_r \mathbf{h}$, $\mathbf{t}_\perp = \mathbf{M}_r \mathbf{t}$
- **Score function**: $f_r(h, t) = -||\mathbf{h}_\perp + \mathbf{r} - \mathbf{t}_\perp||$



3) DistMult

(Yang et al. 2015)

- So far: The scoring function $f_r(h, t)$ is **negative of L1 / L2 distance** in TransE and TransR
- Idea: Use **bilinear** modeling:
Score function: $f_r(h, t) = h \cdot A \cdot t$
 $h, t \in \mathbb{R}^k, A \in \mathbb{R}^{k \times k}$

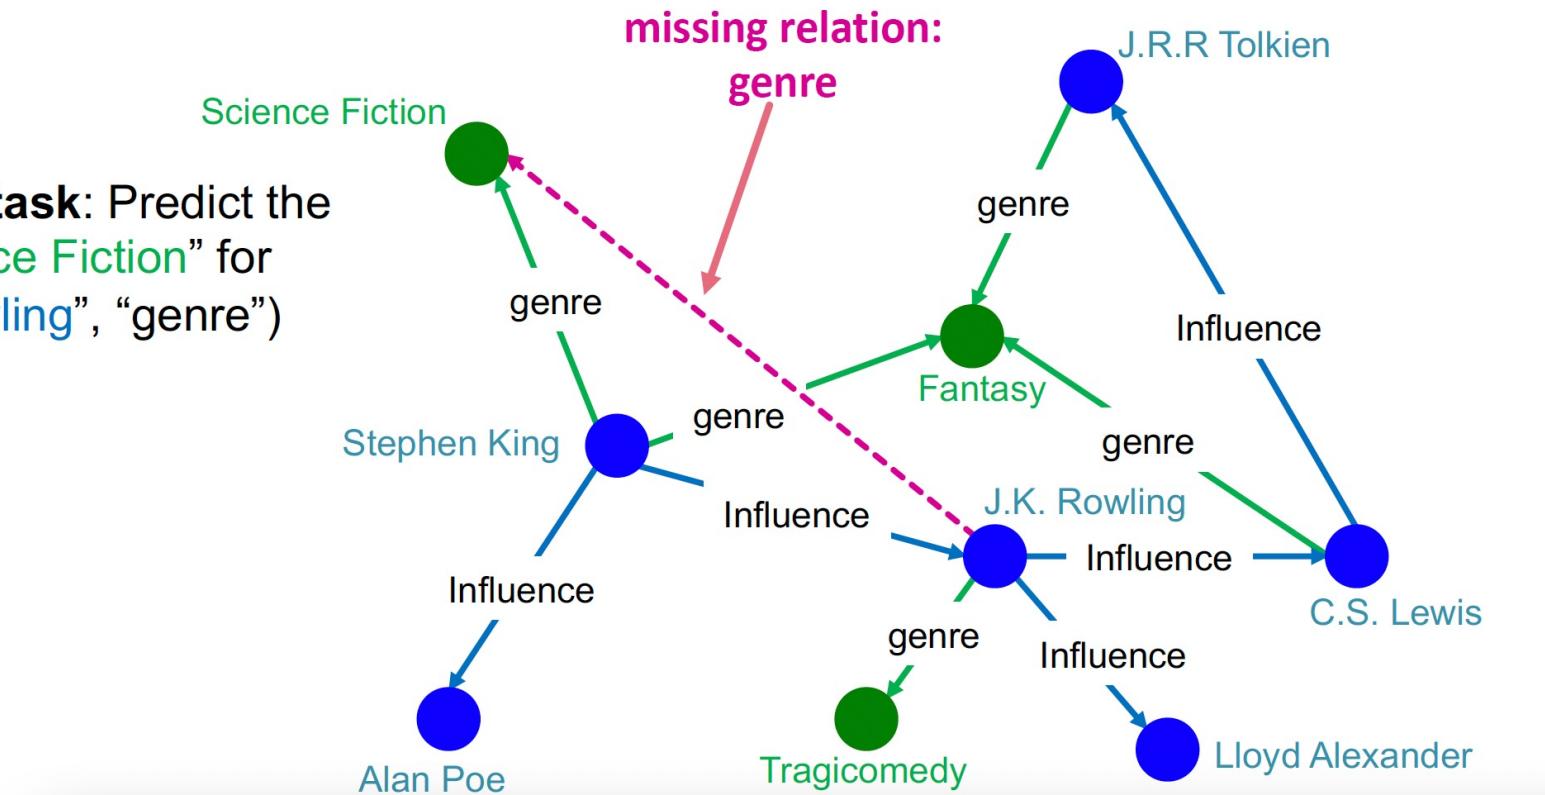
Answering queries (Reasoning) with Knowledge Graphs

Knowledge Graph Completion

Given an enormous KG, can we complete the KG?

- For a given (**head**, **relation**), we predict missing **tails**.

Example task: Predict the tail “**Science Fiction**” for (“**J.K. Rowling**”, “**genre**”)



We can use f_r for this prediction, i.e., is $f_r > t$?

Multi-hop reasoning on KGs

Can we do multi-hop reasoning, i.e., answer complex queries on an incomplete, massive KG?

Query Types	Examples: Natural Language Question, Query
One-hop Queries	What adverse event is caused by Fulvestrant? (e:Fulvestrant, (r:Causes))
Path Queries	What protein is associated with the adverse event caused by Fulvestrant? (e:Fulvestrant, (r:Causes, r:Assoc))
Conjunctive Queries	What is the drug that treats breast cancer and caused headache? ((e:BreastCancer, (r:TreatedBy)), (e:Migraine, (r:CausedBy)))

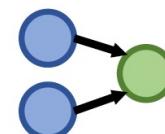
In this lecture, we only focus on answering **queries** on a KG!
The notation will be detailed next.



One-hop Queries



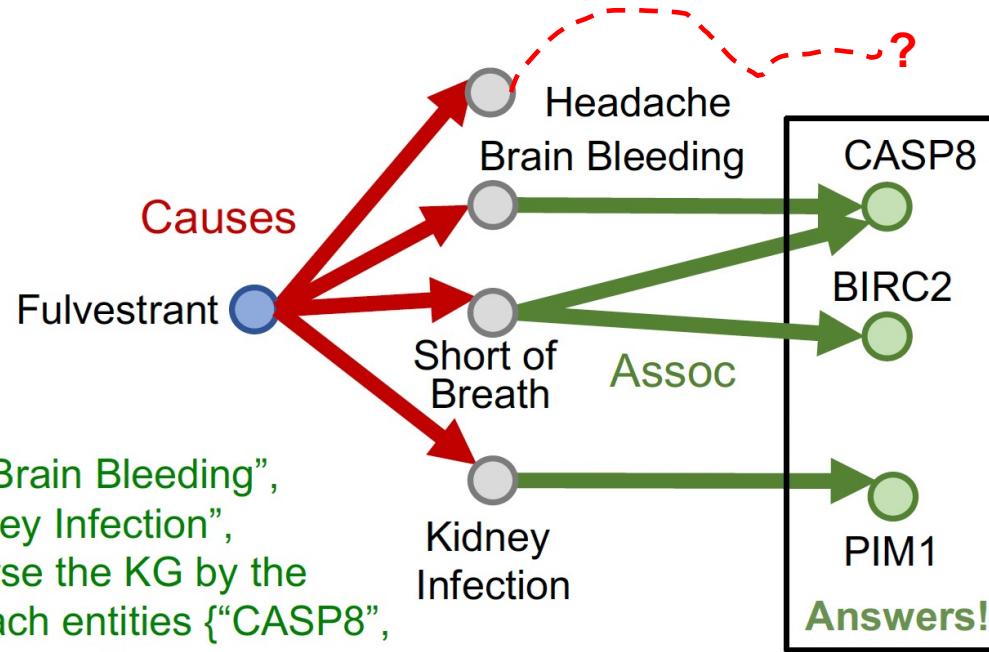
Path Queries



Conjunctive Queries

Traversing the Knowledge Graph

- We answer path queries by traversing the KG:
“What proteins are **associated** with adverse events **caused** by **Fulvestrant**? ”
- Query: (e:Fulvestrant, (r:Causes, r:Assoc))

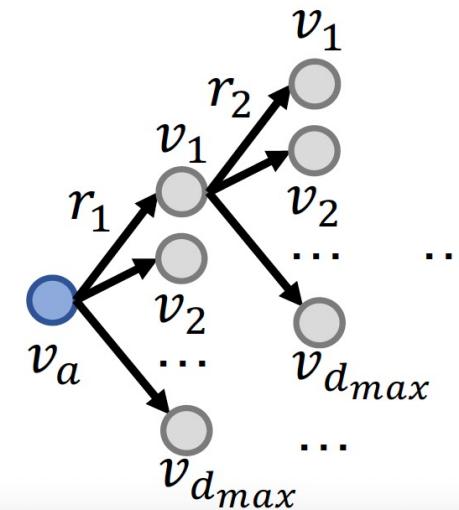


However, knowledge graphs are incomplete!

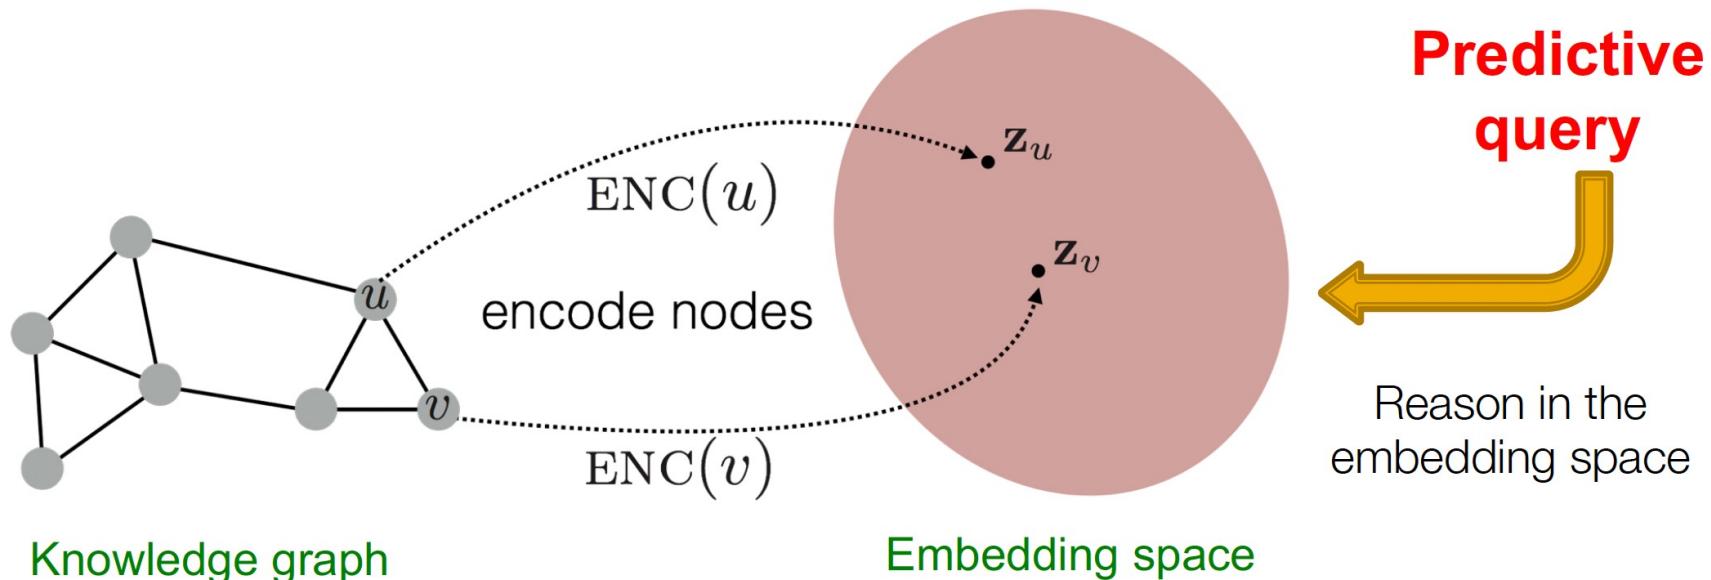
Can KG completion help?

Can we first do KG completion and then traverse the completed (probabilistic) KG?

- **No!** The “completed” KG is a **dense graph**!
 - Most (h, r, t) triples (edge on KG) will have some non-zero probability.
- Time complexity of traversing a dense KG is exponential as a function of the path length L : $O(d_{max}^L)$



Couple of General Ideas



Map queries into embedding space. **Learn to reason in that space**

- Embed query into a single **point** in the Euclidean space: answer nodes are close to the query.
- **Query2Box:** Embed query into a hyper-rectangle (**box**) in the Euclidean space: answer nodes are enclosed in the box.

[[Embedding Logical Queries on Knowledge Graphs](#). Hamilton, et al., NeurIPS 2018]

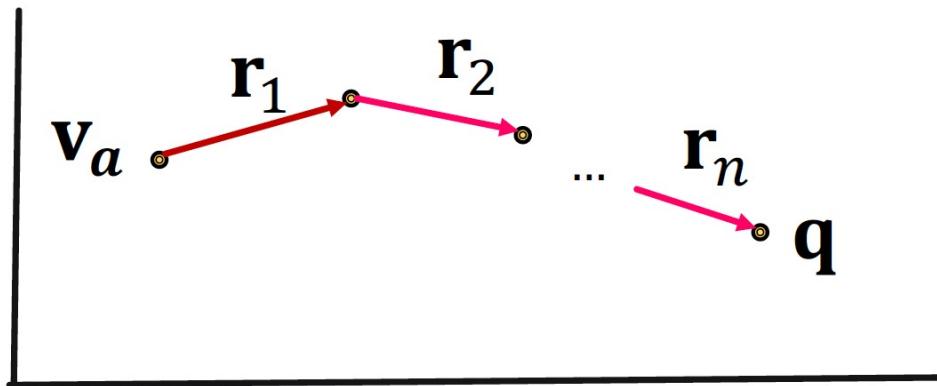
[[Query2box: Reasoning over Knowledge Graphs in Vector Space Using Box Embeddings](#). Ren, et al., ICLR 2020]

An Approach to Query Answering with TransE

- **Key idea: Embed queries!**

- Generalize **TransE** to multi-hop reasoning.

Given a path query $q = (v_a, (r_1, \dots, r_n))$,



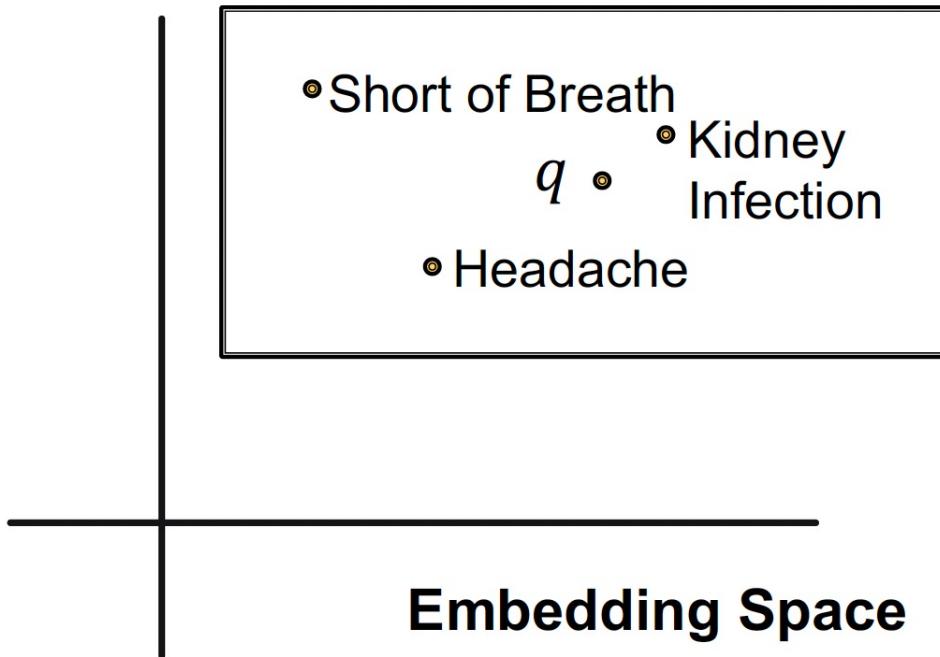
$$q = v_a + r_1 + \cdots + r_n$$

- The embedding process **only involves vector addition, independent of # entities** in the KG!

Another approach: Query2Box

- Embed queries with **hyper-rectangles (boxes)**

$$\mathbf{q} = (\text{Center}(q), \text{Offset}(q))$$



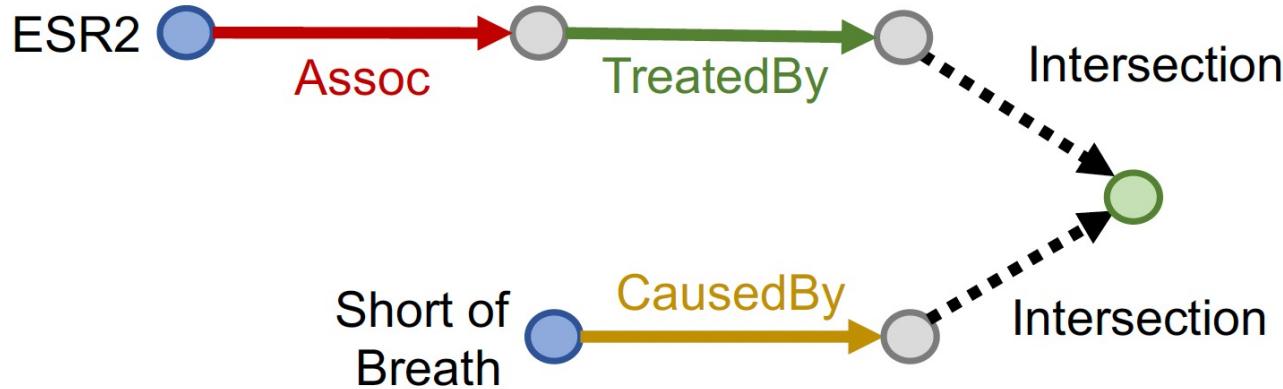
For example, we can embed the adverse events of Fulvestrant with a box that enclose all the answer entities.

Query Answering with Query2Box

How can we answer **more complex queries with logical conjunction operation?**

Query plan:

“What are drugs that cause Short of Breath and treat diseases associated with protein ESR2?”



(1) Each intermediate node represents a set of entities (box)

Reasoning with Query2Box

- **Embed queries in vector space:** “*What are drugs that cause Short of Breath and treat diseases associated with protein ESR2?*”
`((e:ESR2, (r:Assoc, r:TreatedBy)), (e:Short of Breath, (r:CausedBy)))`

Traverse KG from anchor nodes: **ESR2** and **Short of Breath**:

Query plan

Embedding Space



ESR2 •

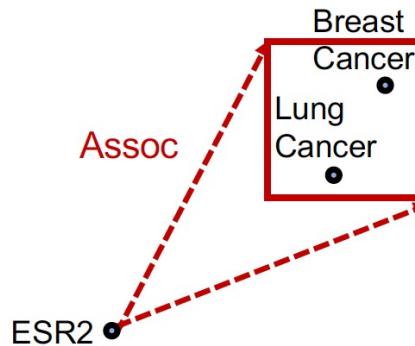
Reasoning with Query2Box

- **Embed queries in vector space:** “*What are drugs that cause Short of Breath and treat diseases associated with protein ESR2?*”
- Traverse KG from anchor nodes: **ESR2** and **Short of Breath**:
- Use **projection operator** again following the query plan.

Query Plan



Embedding Space



Reasoning with Query2Box

“What is the drug that causes Short of Breath and treats disease associated with protein ESR2?”

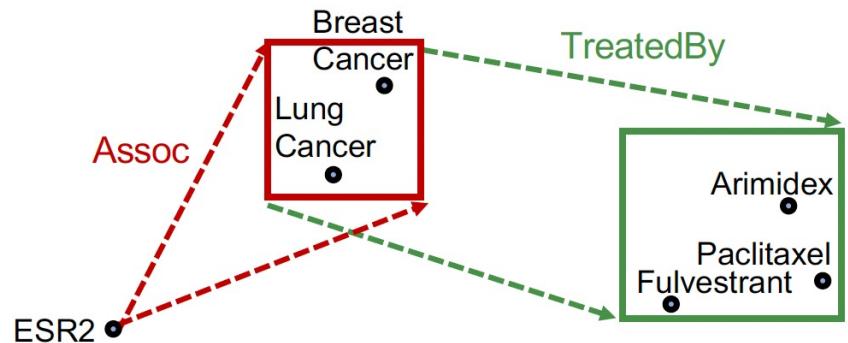
((e:ESR2, (r:Assoc, r:TreatedBy)), (e:Short of Breath, (r:CausedBy)))

- Use **projection operator** again following the query plan.

Query Plan



Embedding Space



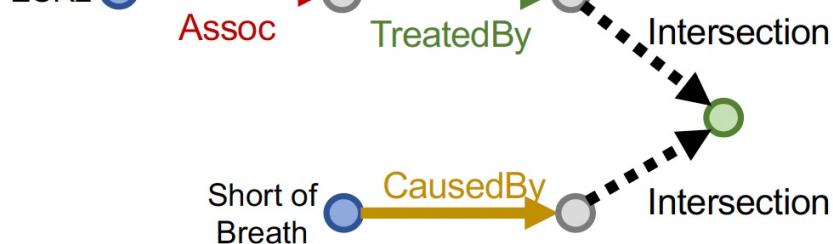
Reasoning with Query2Box

“What is the drug that causes Short of Breath and treats disease associated with protein ESR2?”

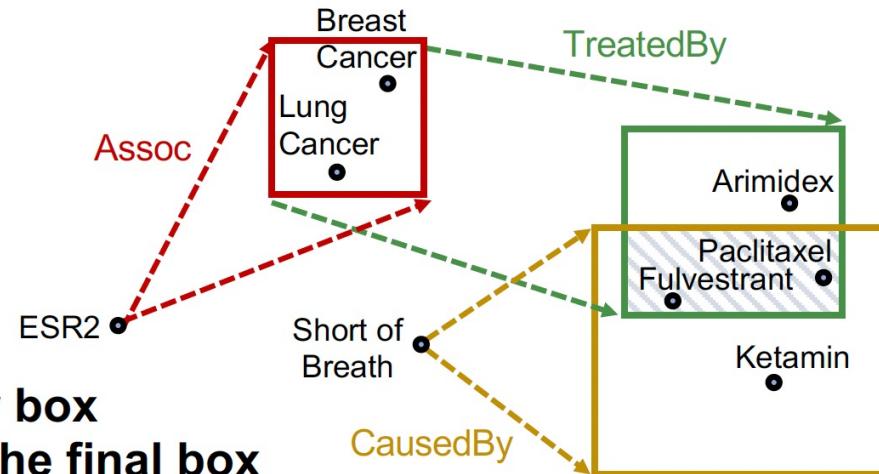
((e:ESR2, (r:Assoc, r:TreatedBy)), (e:Short of Breath, (r:CausedBy)))

■ Use box intersection operator

Query Plan



Embedding Space



**The shadow box
represents the final box
embedding of the query**

Summary

- Introduced KG representation learning models with different embedding spaces and expressiveness - TransE / TransR / DistMult / ...
- Introduced approaches for answering queries on KGs
 - The key idea is to embed queries by navigating the embedding space
 - We embed the query by composing learned operators
 - Embedding of the query is close to its answers in the embedding space

Thanks!